

William Stallings
Computer Organization
and Architecture
7th Edition

Chapter 4
Cache Memory

Consider from key points

- What is the **hierarchical organization** of computer memory?
- What are **the characteristics** of memory hierarchy from top down?
- What is the **cache** for?

Key points

- Computer memory is organized into a hierarchy. (***processor registers** ← **L1 (L2...) cache** ← **main memory** ← **external memory**)
- **Decreasing cost/bit, increasing capacity, slower access time**
- **Locality** of reference makes proper designs of cache improve computer performance

4.1 Computer memory system overview

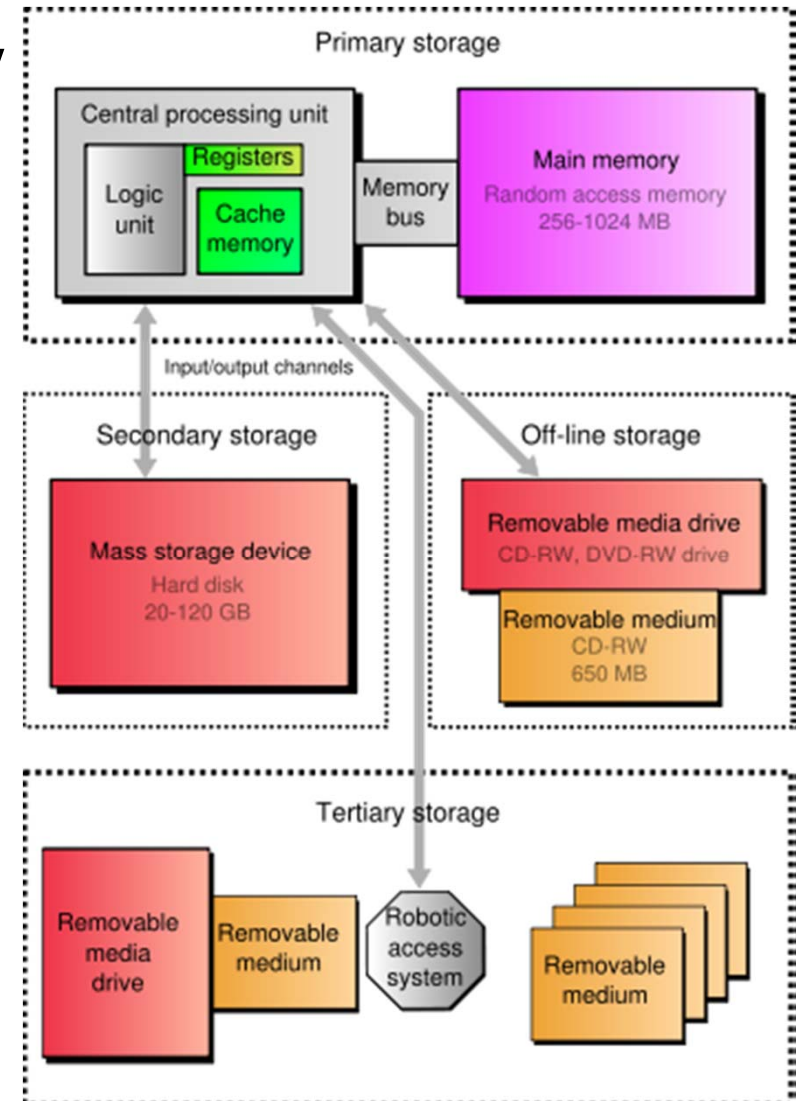
- Characteristics of memory systems
- The memory hierarchy

Characteristics

- Location
- Capacity
- Unit of transfer
- Access method
- Performance
- Physical type
- Physical characteristics
- Organization

Location

- CPU:
 - Register, Control memory
- Internal
 - Cache, Main memory
- External
 - Disk, Tape, CD, DVD



Capacity

- Word size
 - The natural unit of organisation
- Number of words
 - or Bytes
- Capacity = ***Word size*** * ***Number of words***


Addressable unit

- Addressable unit
 - Smallest location which can be uniquely addressed
 - *Word* internally
 - *Cluster* disks

Unit of Transfer


- Internal

- Usually governed by *data bus width* (Byte, Word)



- External

- Usually a *block* which is much larger than a word (Block)



Access Methods (1)

- Sequential
 - Start at the beginning and read through in order
 - Access time depends on *location of data* and *previous location*
 - e.g. tape
- Direct
 - Individual blocks have *unique* address
 - Access is by jumping to vicinity *plus sequential search*
 - Access time depends on *location* and *previous location*
 - e.g. disk

Access Methods (2)

- Random
 - Individual addresses identify locations exactly
 - Access time is independent of location or previous access
 - e.g. RAM
- Associative
 - Data is located by a *comparison* with contents of a portion of the store
 - Access time is independent of location or previous access
 - e.g. cache

Performance

- Access time
 - Time between presenting *the address* and getting the *valid data*
- Memory Cycle time
 - Time may be required for the memory to "*recover*" before next access
 - Cycle time is *access + recovery*
- Transfer Rate R
 - Rate at which data can be moved
 - $1/T$

$$T_N = T_A + N / R$$

Physical Types

- Semiconductor
 - RAM
- Magnetic
 - Disk & Tape
- Optical
 - CD & DVD



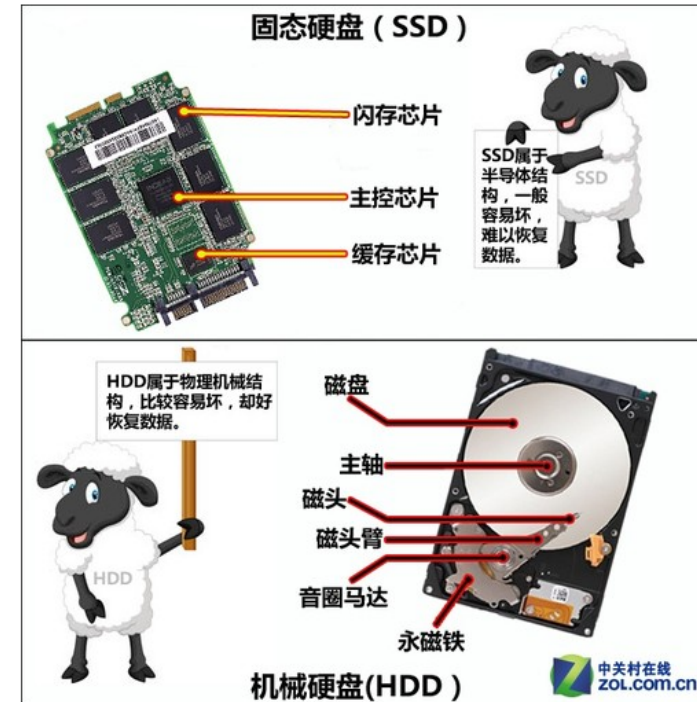
COMPACT
disc
Recordable

Physical Characteristics

- Decay
- Volatility
- Erasable
- Power consumption



脚本之家
WWW.JB51.NET




Memory design

- How much?
 - Capacity
- How fast?
 - Time is money
- How expensive?
- Dilemma:
 - Faster access time, greater cost per bit
 - Greater capacity, smaller cost per bit
 - Greater capacity, slower access time

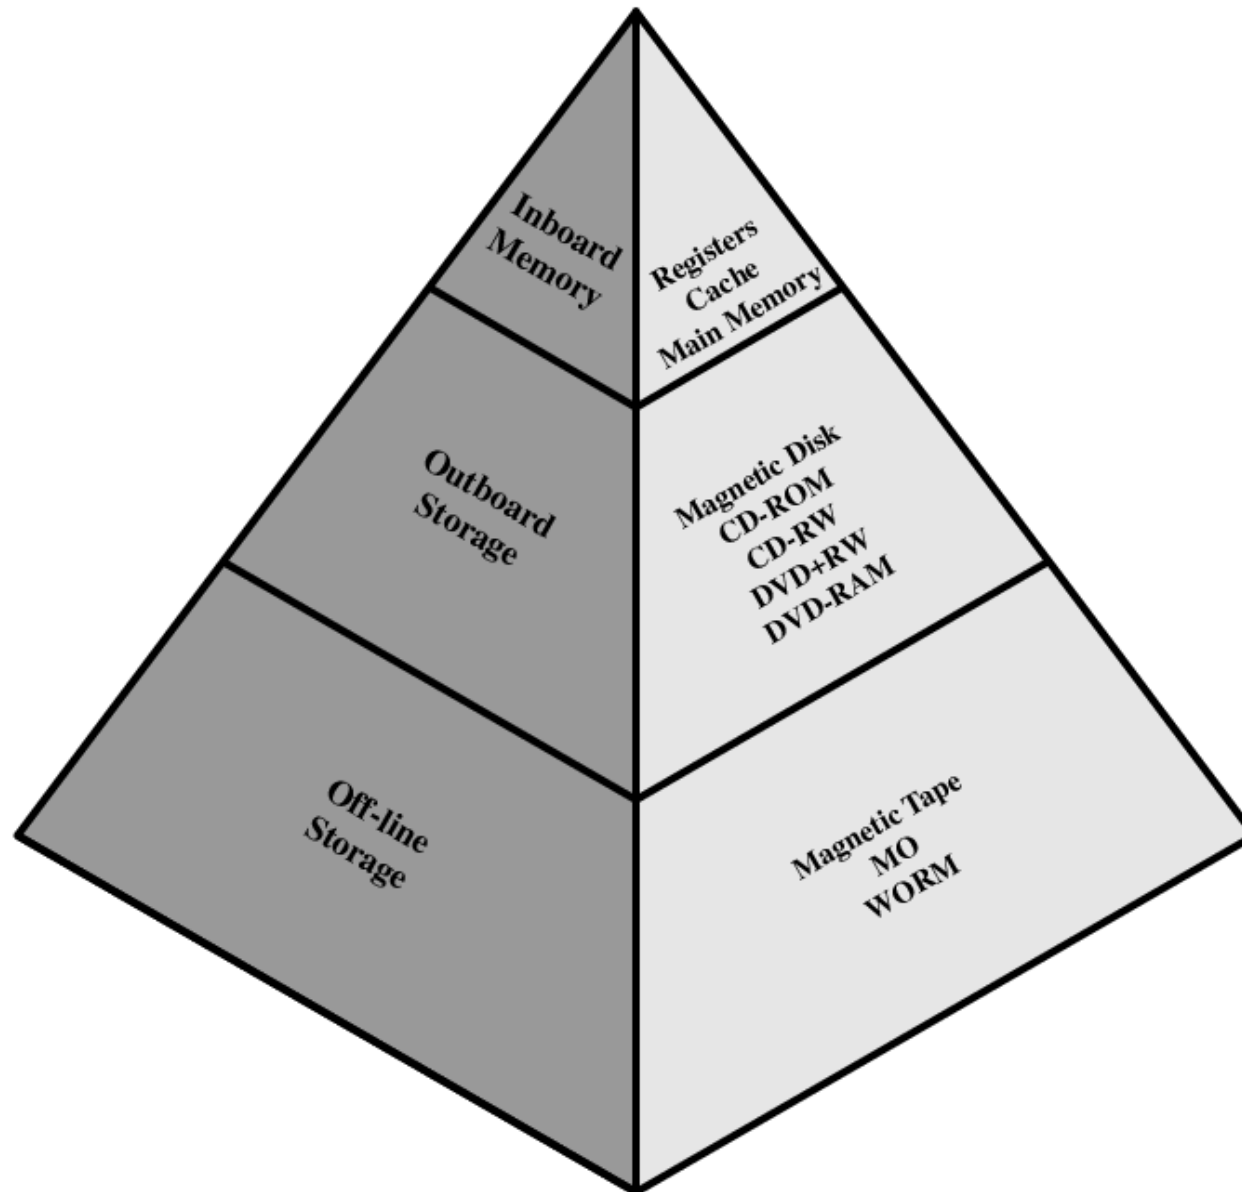


Memory Hierarchy

- Registers
 - In CPU
- Internal or Main memory
 - May include one or more levels of cache
 - "RAM"
- External memory
 - Backing store

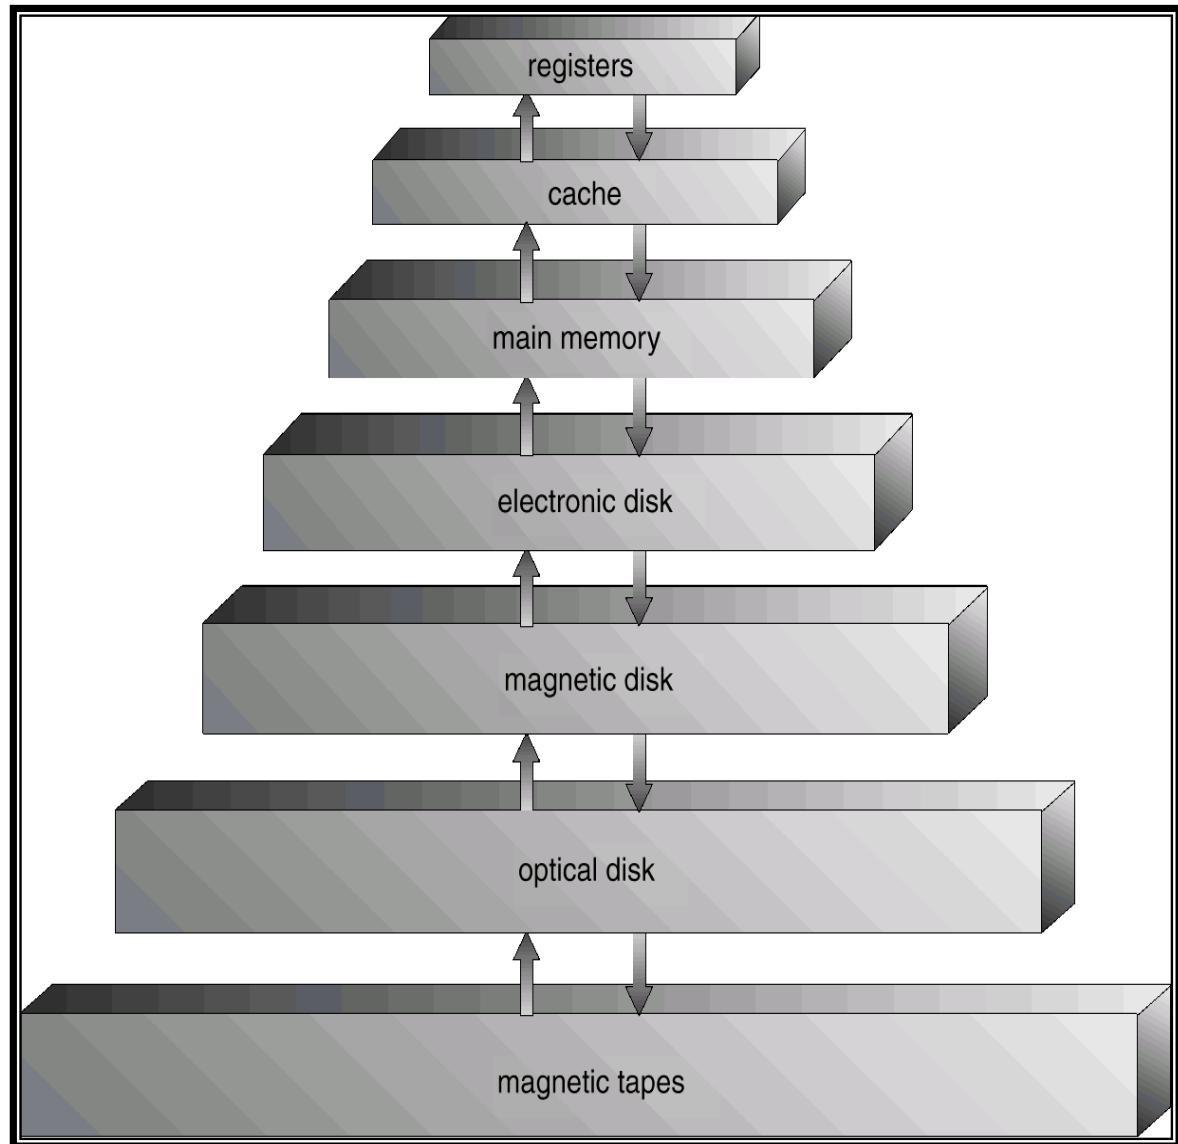


Memory Hierarchy - Diagram




Hierarchy List

- Registers
- L1 Cache
- L2 Cache
- Main memory
- *Disk cache*
- Disk
- Optical
- Tape



Performance of two-level memory




$$T = T_1 H + (T_1 + T_2)(1 - H) \quad (4-1)$$

Locality of reference (Appendix 4A)


- Behaviour of programs
 - Sequential execution
 - Few procedures (Fig. 4.16)
 - Iterative constructs
 - Arrays or sequences of records
- Spatial locality and Temporal locality

4.2 Cache memory principles


- Small amount of fast memory
- Sits between normal main memory and CPU
- May be located on CPU chip or module



Cache/Main Memory Structure



(a) Cache

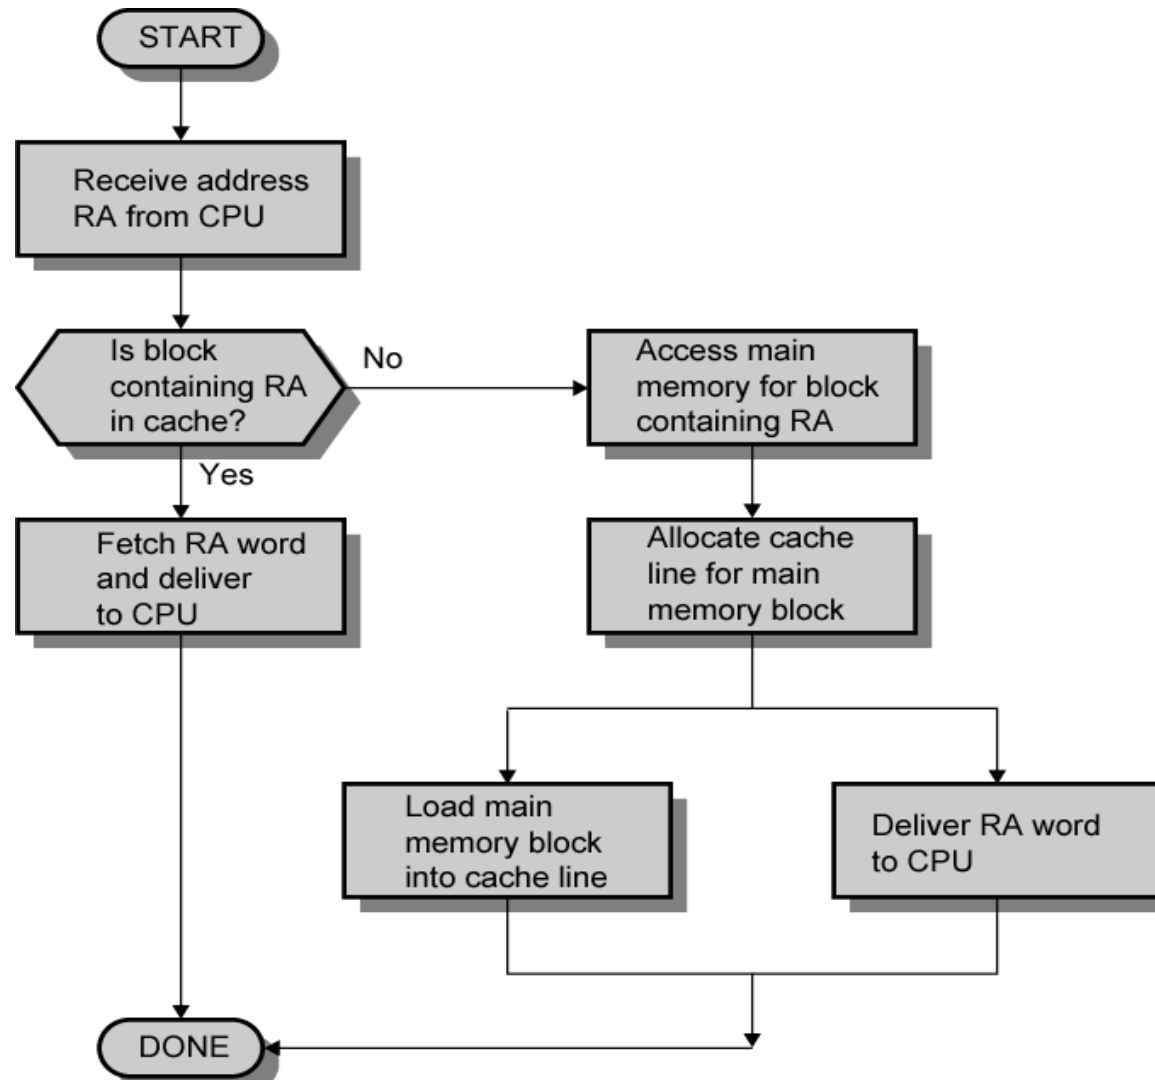


(b) Main memory

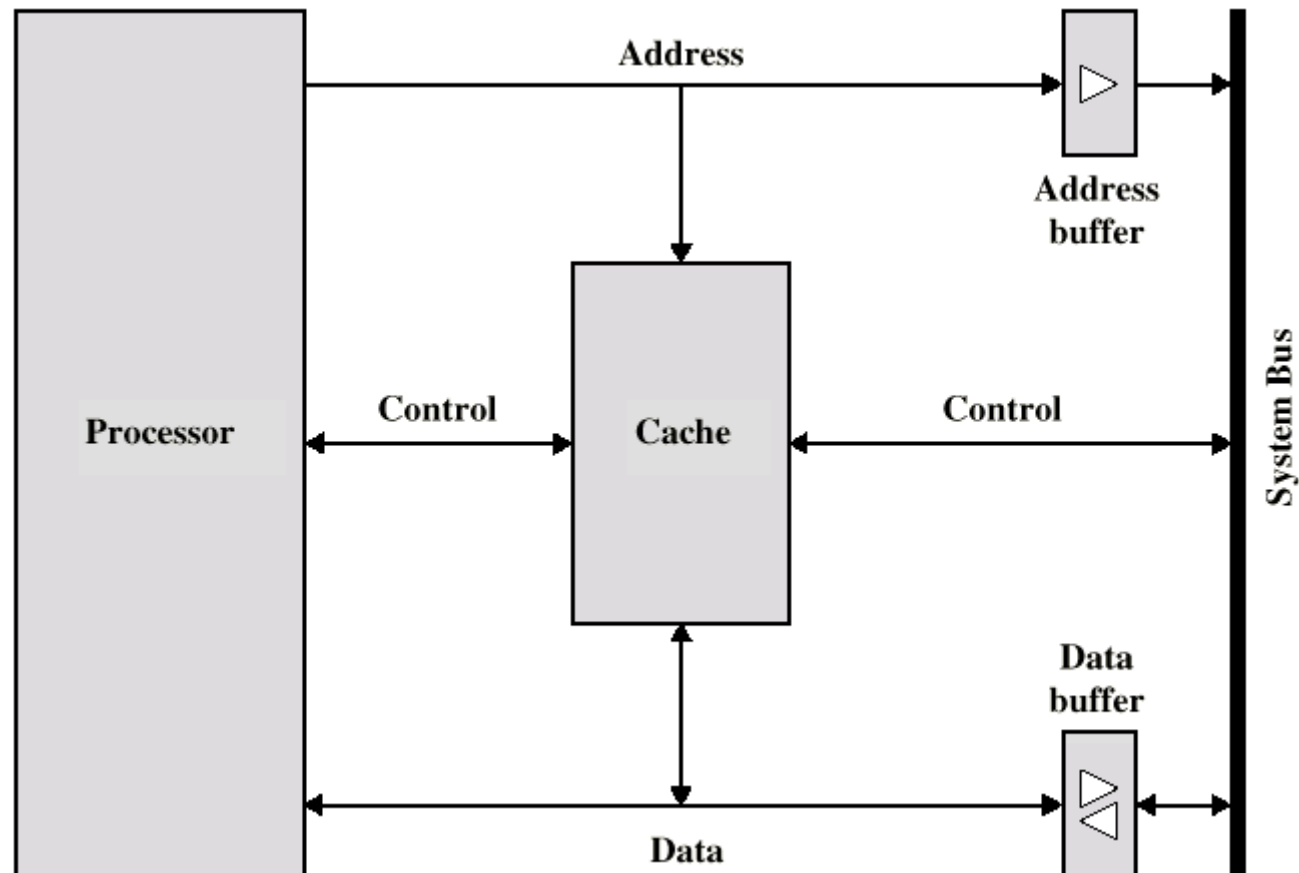
Cache operation – overview

- CPU requests contents of **memory location**
- Check cache for this data
- If present, get from cache (**hit**)
- If not present (**miss**), read required **block** from main memory to cache Then deliver from cache to CPU
- Cache includes **tags** to identify which block of main memory is in each cache slot

Cache Read Operation - Flowchart



Typical Cache Organization



4.3 Elements of Cache Design

- Size
- Mapping Function
- Replacement Algorithm
- Write Policy
- Block Size
- Number of Caches

Size does matter

- Cost
 - More cache is expensive
- Speed
 - More cache is faster
 - Checking cache for data takes time

Comparison of Cache Sizes

Processor	Type	Year of Introduction	L1 cache ^a	L2 cache	L3 cache
IBM 360/85	Mainframe	1968	16 to 32 KB	—	—
PDP-11/70	Minicomputer	1975	1 KB	—	—
VAX 11/780	Minicomputer	1978	16 KB	—	—
IBM 3033	Mainframe	1978	64 KB	—	—
IBM 3090	Mainframe	1985	128 to 256 KB	—	—
Intel 80486	PC	1989	8 KB	—	—
Pentium	PC	1993	8 KB/8 KB	256 to 512 KB	—
PowerPC 601	PC	1993	32 KB	—	—
PowerPC 620	PC	1996	32 KB/32 KB	—	—
PowerPC G4	PC/server	1999	32 KB/32 KB	256 KB to 1 MB	2 MB
IBM S/390 G4	Mainframe	1997	32 KB	256 KB	2 MB
IBM S/390 G6	Mainframe	1999	256 KB	8 MB	—
Pentium 4	PC/server	2000	8 KB/8 KB	256 KB	—
IBM SP	High-end server/ supercomputer	2000	64 KB/32 KB	8 MB	—
CRAY MTA ^b	Supercomputer	2000	8 KB	2 MB	—
Itanium	PC/server	2001	16 KB/16 KB	96 KB	4 MB
SGI Origin 2001	High-end server	2001	32 KB/32 KB	4 MB	—
Itanium 2	PC/server	2002	32 KB	256 KB	6 MB
IBM POWER5	High-end server	2003	64 KB	1.9 MB	36 MB
CRAY XD-1	Supercomputer	2004	64 KB/64 KB	1MB	—

^a Two values separated by a slash refer to instruction and data caches

^b Both caches are instruction only; no data caches

Mapping Function

- Direct
- Associative
- Set associative

Mapping example

- Cache of 64kByte
- Cache block of 4 bytes
 - i.e. cache is 16k (2^{14}) lines of 4 bytes
- 16MBytes main memory
- 24 bit address
 - i.e. $2^{24} = 16\text{M}$

Direct Mapping

- Each block of main memory maps to only one cache line
 - i.e. if a block is in cache, it must be in **one specific place**
- Address is in two parts
- **Least Significant w bits** identify unique word or byte
- **Most Significant s bits** specify one memory block
- The MSBs are split into a cache line field r and a tag of $s-r$ (most significant)

Direct Mapping Address Structure

Tag s-r	Line or Slot r	Word w
8	14	2

- 24 bit address (main memory)
- 2 bit word identifier (4 byte block)
- 22 bit block identifier
 - 8 bit tag (=22-14)
 - 14 bit slot or line (Cache line number)
- No two blocks in the same line have the same Tag field
- Check contents of cache by finding line and checking Tag

Direct Mapping function

$$i = j \bmod m$$

i = cache set number

j = main memory block number

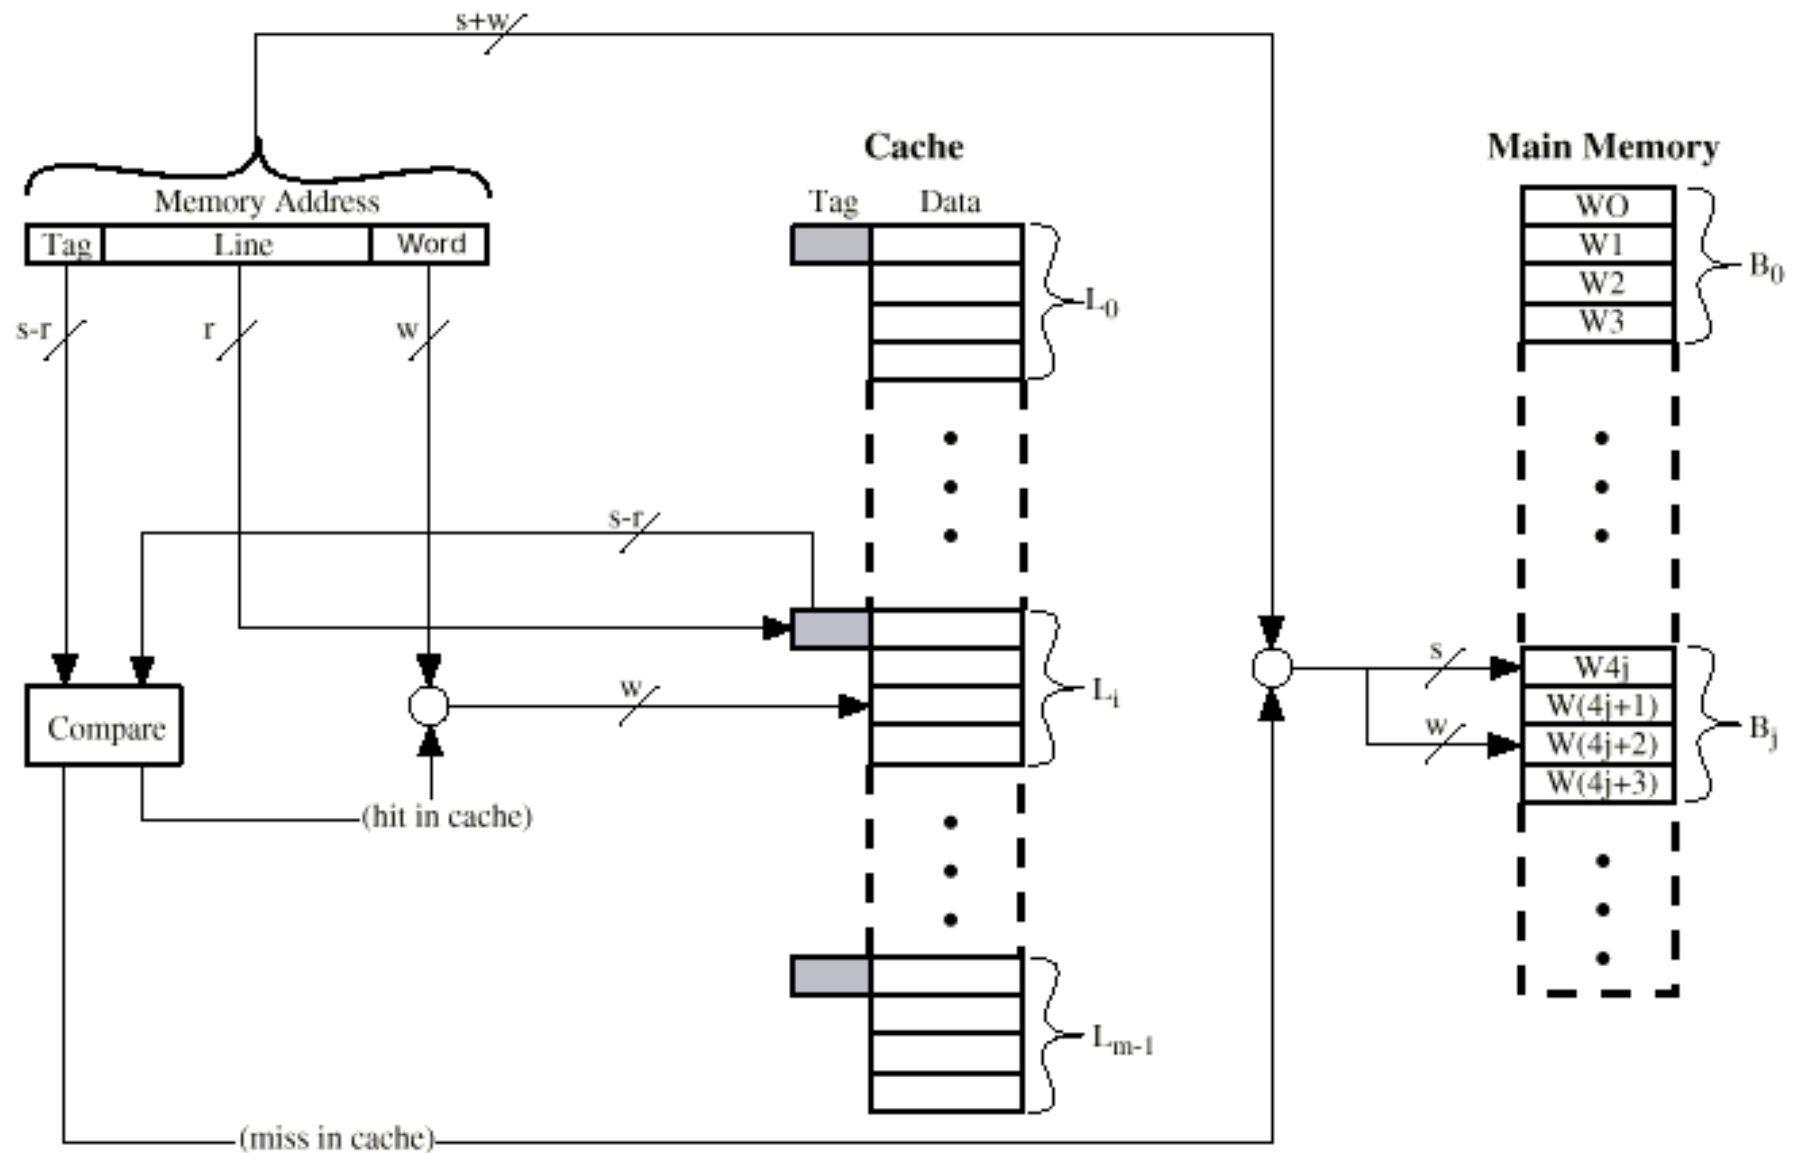
m = number of lines in the cache

Direct Mapping

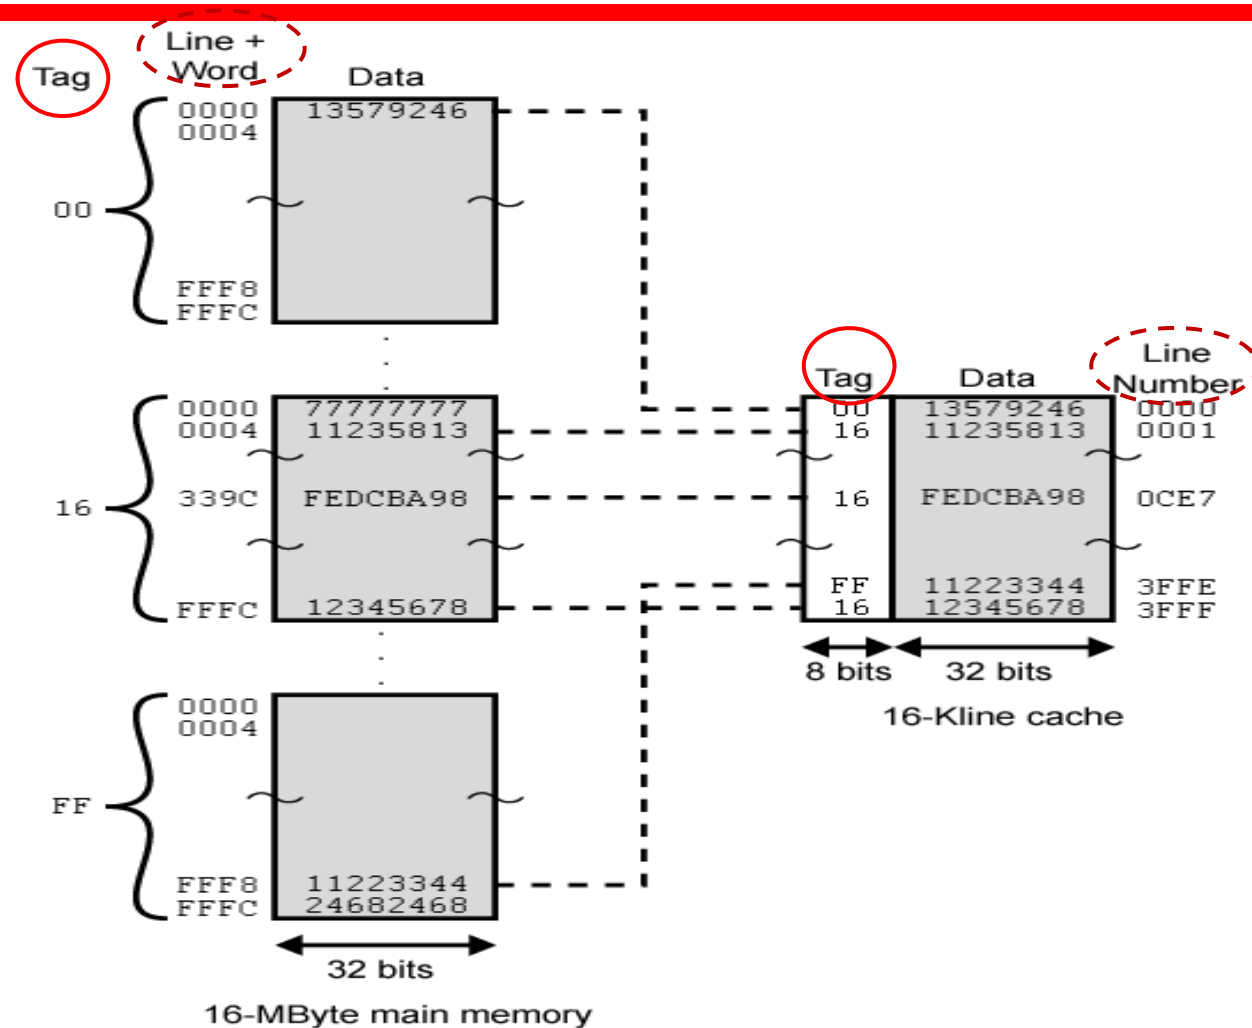
Cache Line Table

• Cache line	Main Memory blocks held
• 0	0, m, 2m, ...2s-m
• 1	1, m+1, 2m+1...2s-m+1
• m-1	m-1, 2m-1, 3m-1...2s-1

Direct Mapping Cache Organization



Direct Mapping Example



	Tag	Line	Word
Main memory address =	8	14	2

Direct Mapping Summary

- Address length = $(s + w)$ bits
- Number of addressable units = 2^{s+w} words or bytes
- Block size = line size = 2^w words or bytes
- Number of blocks in main memory = $2^{s+w} / 2^w = 2^s$
- Number of lines in cache = $m = 2^r$
- Size of tag = $(s - r)$ bits

Advantages and disadvantages of Direct Mapping

Adv.

- Simple
- Inexpensive

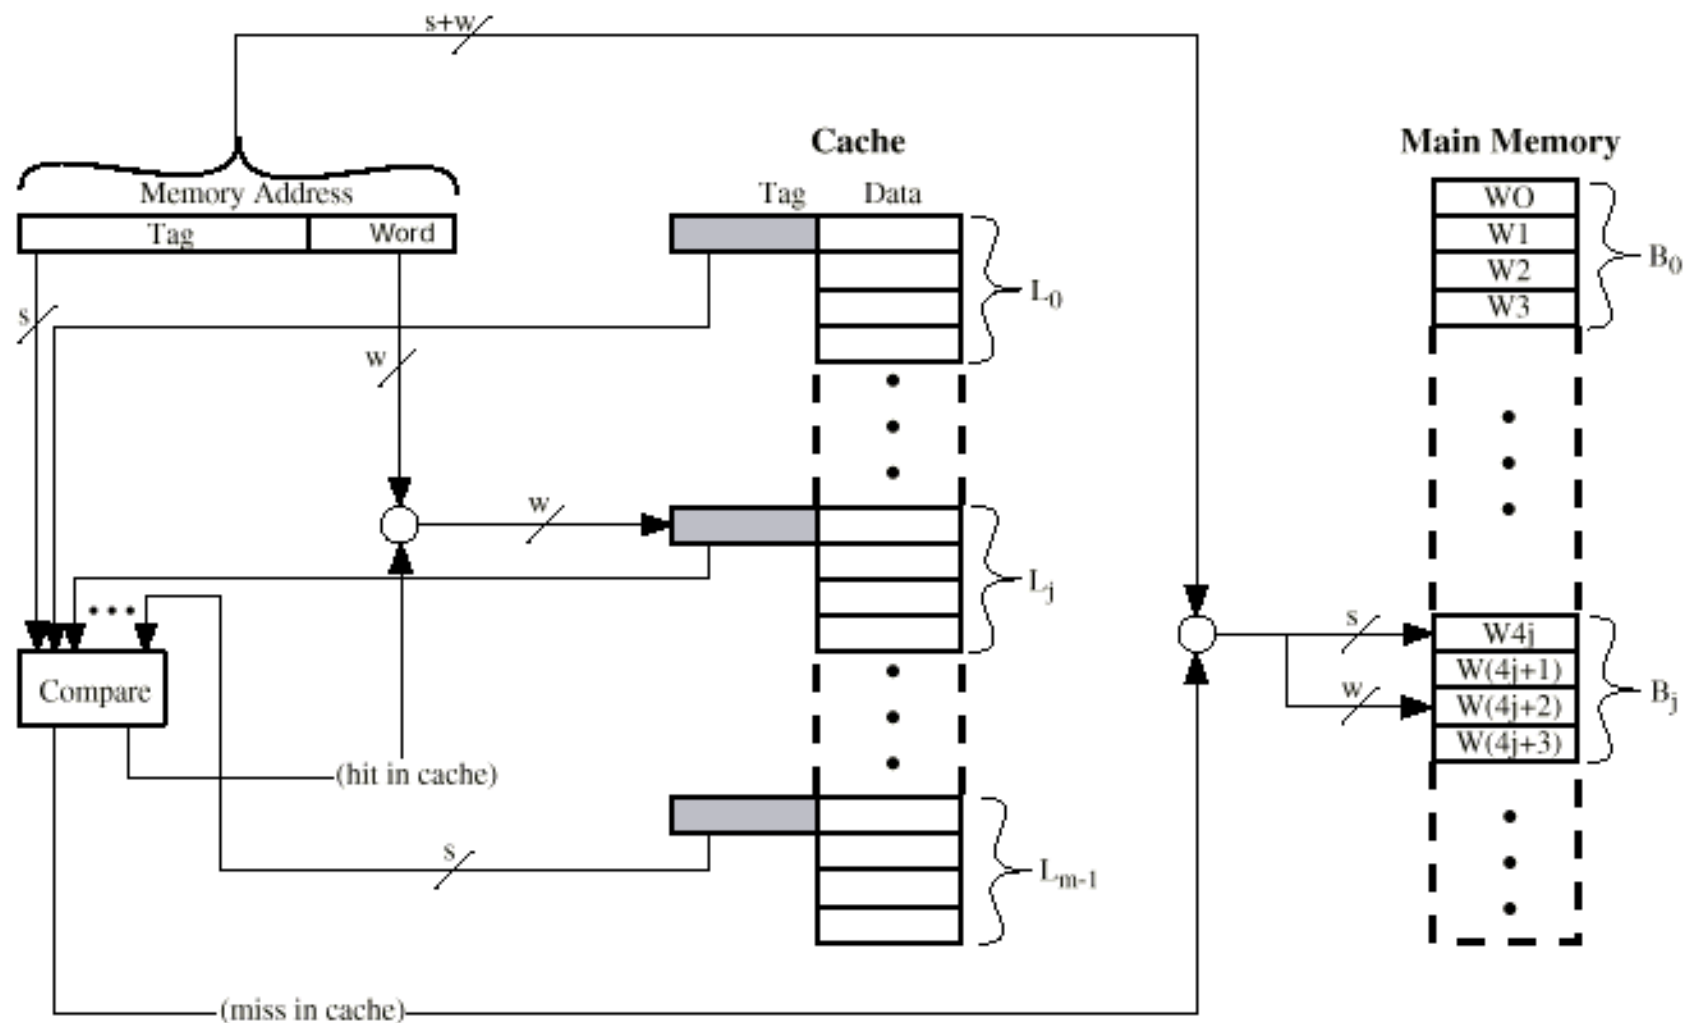
Disadv.

- Fixed location for given block
 - If a program accesses 2 blocks that map to the same line repeatedly, cache misses are very high

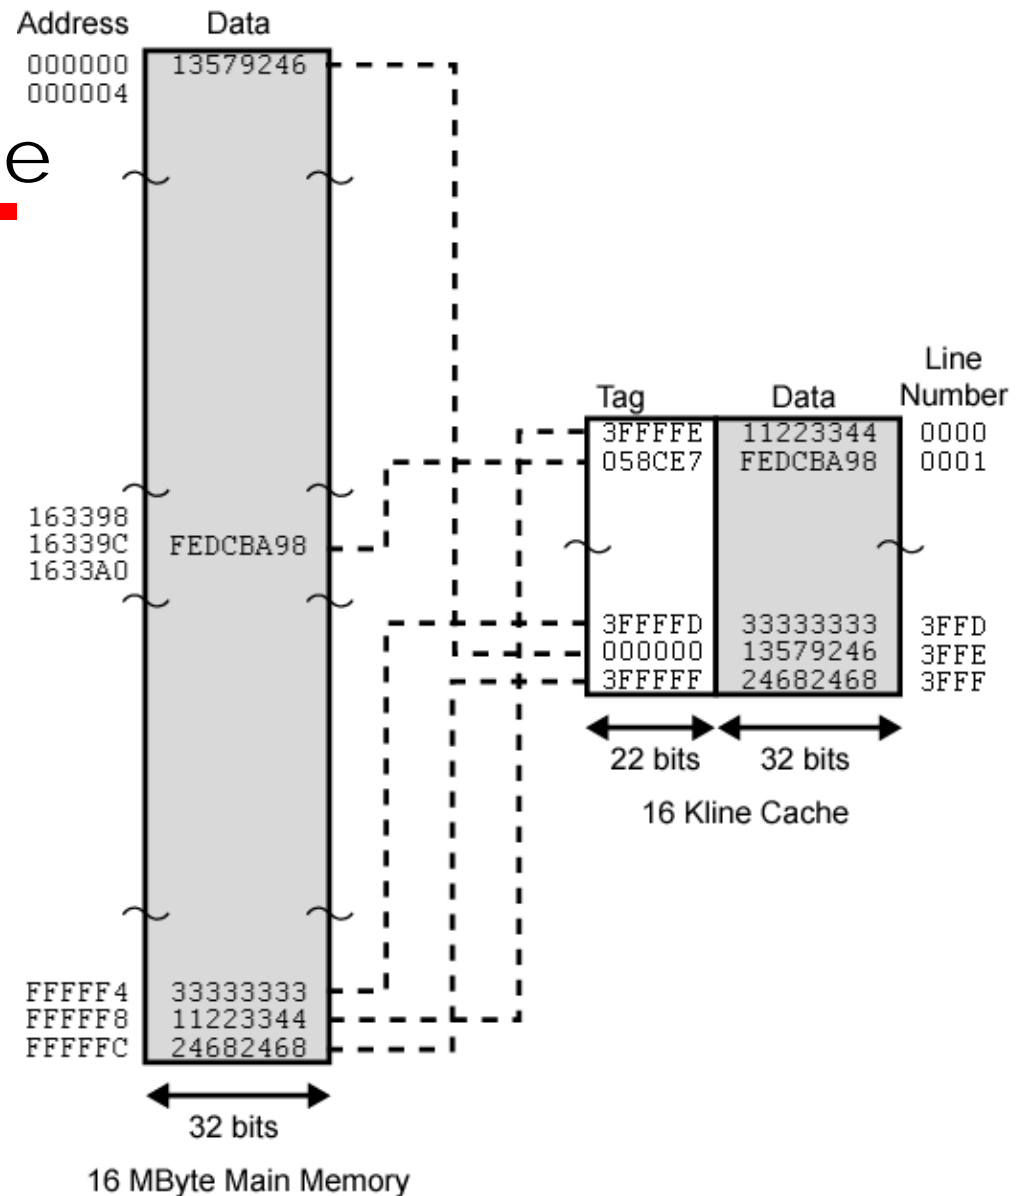
Associative Mapping

- A main memory block can load into **any** line of cache
- Memory address is interpreted as **tag and word**
- Tag *uniquely* identifies block of memory
- *Every* line's tag is examined for a match
- Cache searching gets **expensive**

Fully Associative Cache Organization



Associative Mapping Example



Main Memory Address =

Tag	Word
22	2

Associative Mapping

Address Structure

Tag 22 bit	Word 2 bit
------------	---------------

- 22 bit tag stored with each 32 bit block of data (1 by 1)
- Compare tag field with tag entry in cache to check for hit
- Least significant 2 bits of address identify which 8 bit byte is required from 32 bit data block

Associative Mapping Summary

- Address length = $(s + w)$ bits
- Number of addressable units = 2^{s+w} words or bytes
- Block size = line size = 2^w words or bytes
- Number of blocks in main memory = $2^{s+w} / 2^w = 2^s$
- Number of line in cache associated to main memory = **undetermined**
- Size of tag = s bits

Set Associative Mapping

- Cache is divided into a number of sets
- Each set contains a number of lines
 - Lines -> **sets** -> cache
- A given block maps to any **line** in a given set
 - e.g. Block B can be in any line of set i
- e.g. 2 lines per set
 - 2 way associative mapping
 - A given block can be in one of 2 lines in only one set

$$m = v \times k$$

$$i = j \bmod v$$

i = cache set No.

j = main memory block No.

m = number of lines in the cache
(Cache size in lines)

v = number of sets in the cache
(Cache size in sets)

Set Associative Mapping

Address Structure

Tag 9 bit	Set 13 bit	Word 2 bit
-----------	------------	------------

- Use set field to **determine** cache set to look in
- **Compare** tag field to see if we have a hit
- e.g

—Address Tag Data Set number

—1FF 7FFC 1FF 12345678 1FFF

—001 7FFC 001 11223344 1FFF

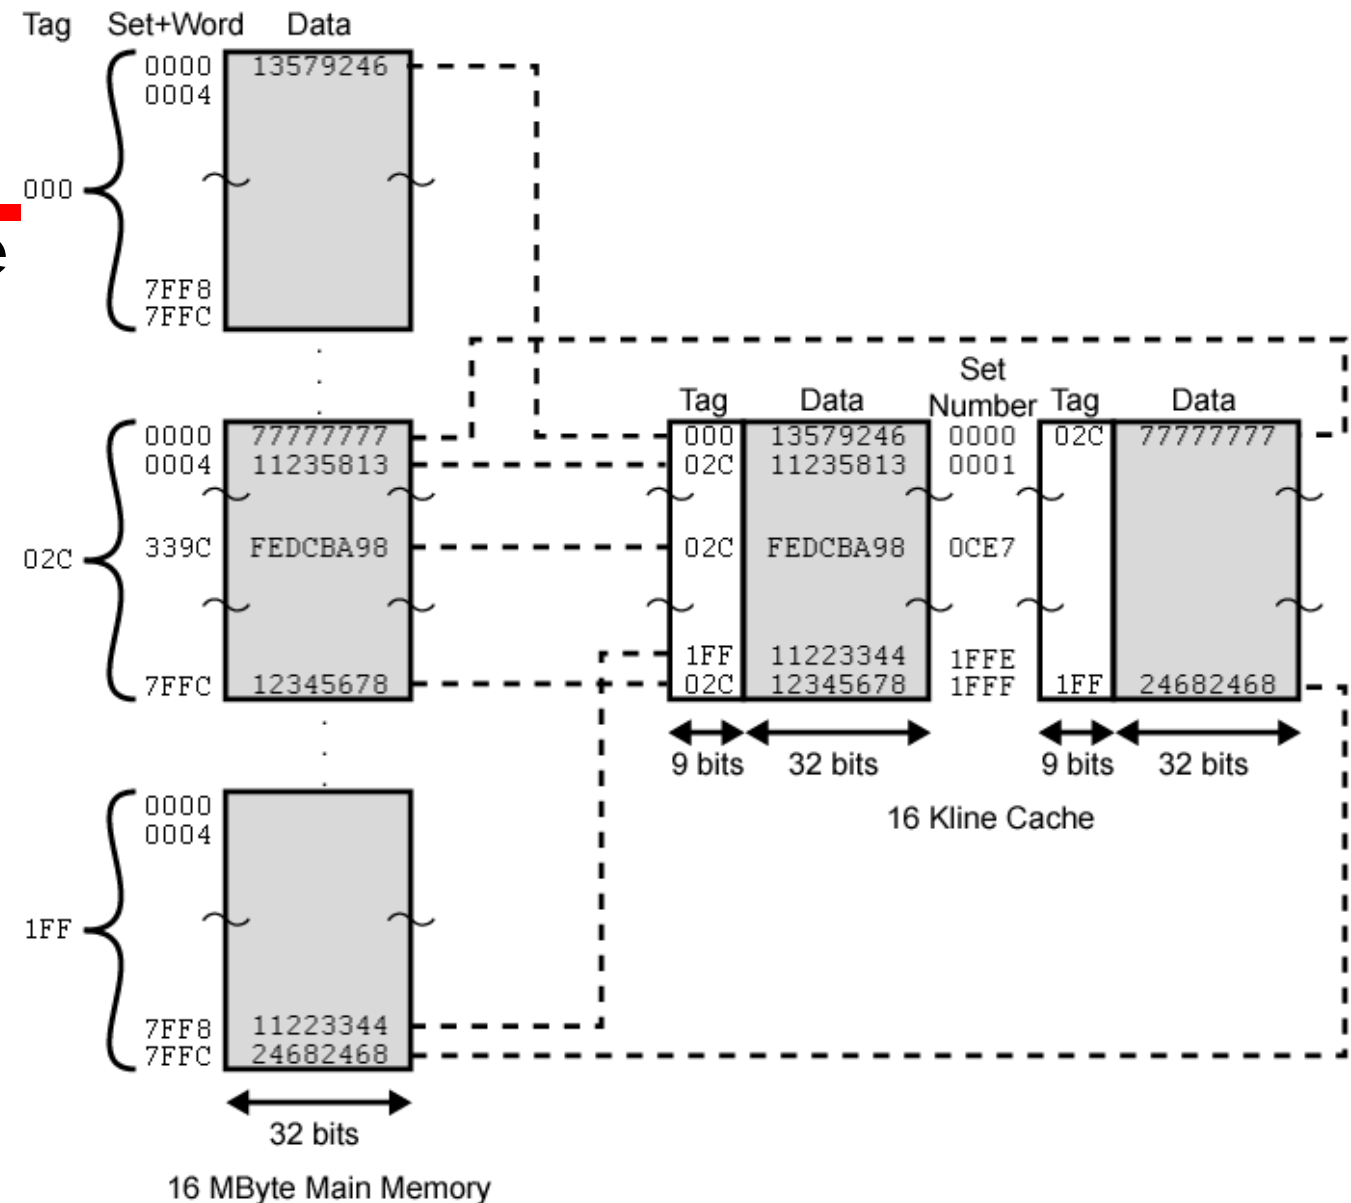
—Set No "1FFF": 0 0000 0001 111 1111 1111 1100

Set Associative Mapping

Example

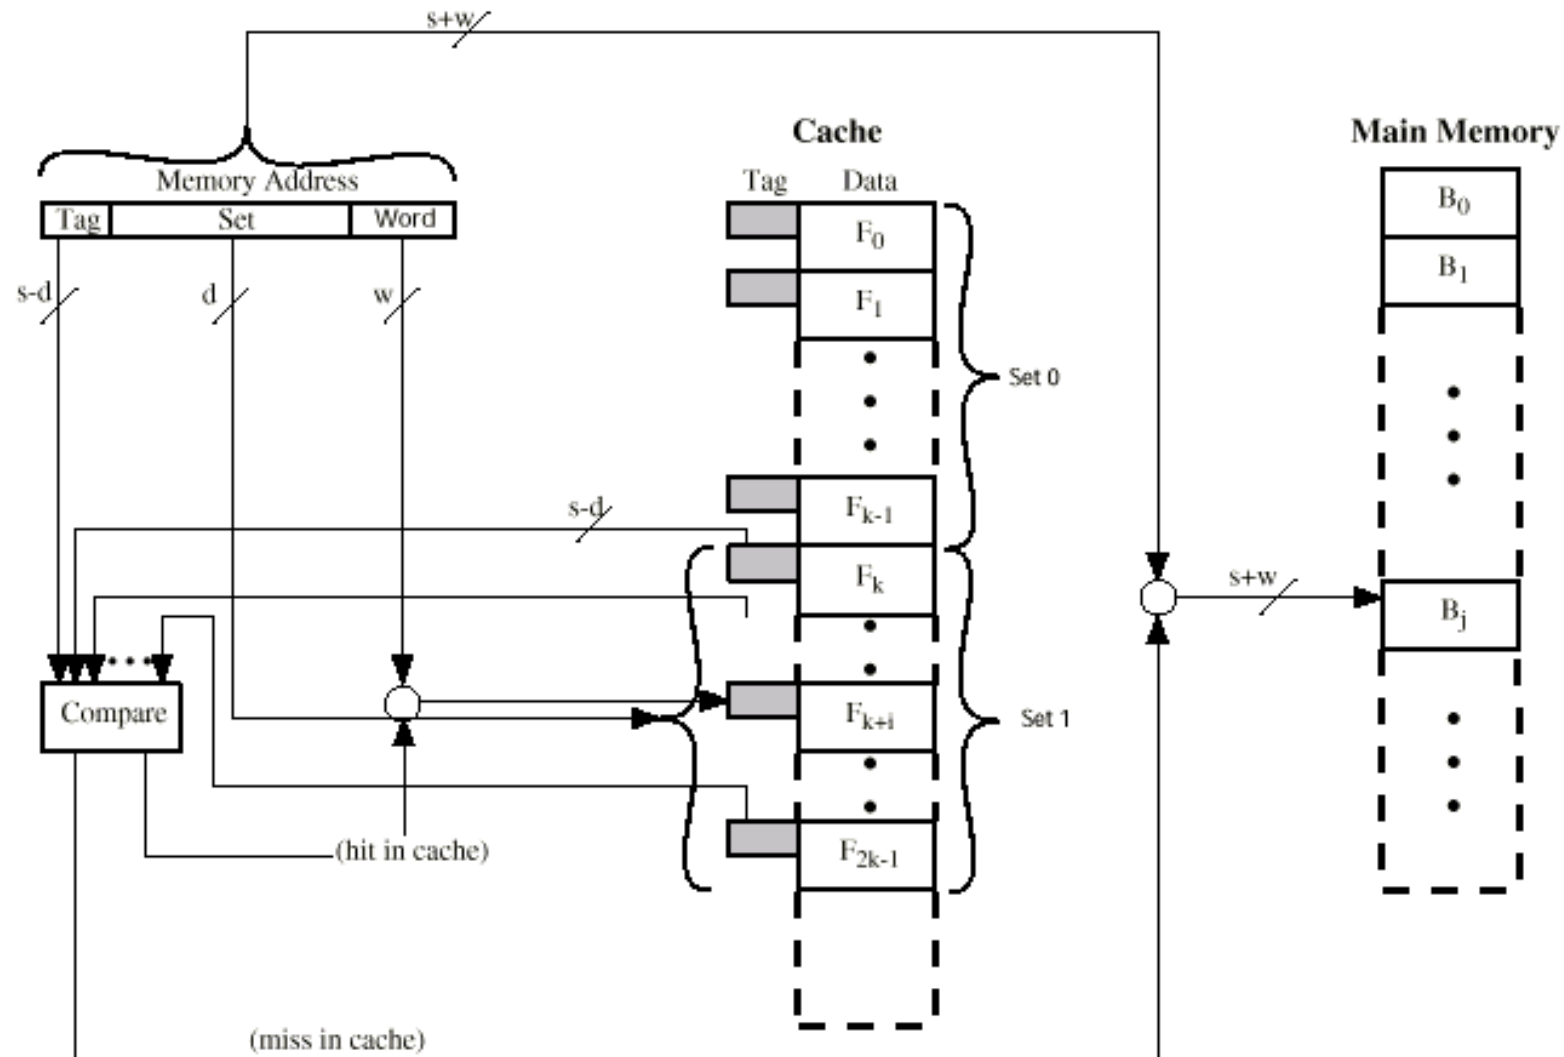
- 13 bit set number
- 2 lines per set
- Block number in main memory is modulo 2^{13}
- 002000, 00A000, 012000, 01A000,...
FF2000, FFA000 ... map to same set 0800

Two Way Set Associative Mapping Example



	Tag	Set	Word
Main Memory Address =	9	13	2

K-Way Set Associative Cache Organization



Set Associative Mapping Summary

- Address length = $(s + w)$ bits
- Number of addressable units = 2^{s+w} words or bytes
- Block size = line size = 2^w words or bytes
- Number of blocks in main memory = 2^s
- Number of lines in one set = k
- Number of sets = $v = 2^d$
- Number of lines in cache = $k v = k * 2^d$
- Size of tag = $(s - d)$ bits

Relationships of mappings

- Set Associative Mapping \rightarrow Direct mapping when $v=m$ and $k=1$
- Set Associative Mapping \rightarrow Associative Mapping when $v=1$ and $k=m$
- $K=2$ is the most common set associative organization \Rightarrow hit improvement
- $K=4$ makes a modest additional improvement

Replacement Algorithms (1)

Direct mapping

- No choice
- Each block only maps to one line
- Replace that line

Replacement Algorithms (2)

Associative & Set Associative

- Least Recently used (LRU)
 - e.g. in 2 way set associative, Which of the 2 block is lru?
- First in first out (FIFO)
 - replace block that has been in cache longest
- Least frequently used
 - replace block which has had fewest hits
- Random
 - Slightly inferior performance to usage based algorithms

Write Policy

- Must not overwrite a cache block unless main memory is **up to date**
- >Problems:
- Multiple CPUs may have individual caches
 - I/O may address main memory directly

Write through

- All writes go to main memory as well as cache
- Multiple CPUs can monitor main memory traffic to keep local (to CPU) cache up to date

->Disadv:

- Lots of traffic
- Slows down writes

Write back

- Updates initially **made in cache only**
- Update bit for cache slot is **set** when update occurs
- If block is to be replaced, write to main memory only if update bit is set
- I/O must access main memory through cache

-> keep single cache & main memory coherent

Cache coherency

- Bus watching with write through
 - Hardware transparency
 - Noncacheable memory
- > keep multiple cache & main memory coherent

Line size

- As the block size increase, hit ratio increase first, then decrease
- 8-32 bytes seems reasonably
- 64 and 128 byte cache line sizes are most frequently used for HPC

Number of caches

- Multilevel caches
 - On chip cache speed up execution
 - Off chip cache is still desirable
- Unified versus split caches
 - Unified cache: only one cache, balance the load between instruction and data fetches,
 - Split cache: One dedicated to instructions and one dedicated to data, suitable for pipeline execution

Pentium 4 Cache

- 80386 – no on chip cache
- 80486 – 8k using 16 byte lines and four way set associative organization
- Pentium (all versions) – two on chip L1 caches
 - Data & instructions
- Pentium 4
 - L1 data/instruction caches
 - 8k bytes
 - 64 byte lines
 - four way set associative
 - L2 cache
 - Feeding both L1 caches
 - 256k
 - 128 byte lines
 - 8 way set associative
 - L3 on chip cache

Intel Cache Evolution

Problem	Solution	Processor on which feature first appears
External memory slower than the system bus.	Add external cache using faster memory technology.	386
Increased processor speed results in external bus becoming a bottleneck for cache access.	Move external cache on-chip, operating at the same speed as the processor.	486
Internal cache is rather small, due to limited space on chip	Add external L2 cache using faster technology than main memory	486
Contention occurs when both the Instruction Prefetcher and the Execution Unit simultaneously require access to the cache. In that case, the Prefetcher is stalled while the Execution Unit's data access takes place.	Create separate data and instruction caches.	Pentium
Increased processor speed results in external bus becoming a bottleneck for L2 cache access.	Create separate back-side bus that runs at higher speed than the main (front-side) external bus. The BSB is dedicated to the L2 cache.	Pentium Pro
	Move L2 cache on to the processor chip.	Pentium II
Some applications deal with massive databases and must have rapid access to large amounts of data. The on-chip caches are too small.	Add external L3 cache.	Pentium III
	Move L3 cache on-chip.	Pentium 4

Pentium 4 Core Processor

- Fetch/Decode Unit
 - Fetches instructions from L2 cache
 - Decode into micro-ops
 - Store micro-ops in L1 cache
- Out of order execution logic
 - Schedules micro-ops
 - Based on data dependence and resources
 - May speculatively execute
- Execution units
 - Execute micro-ops
 - Data from L1 cache
 - Results in registers
- Memory subsystem
 - L2 cache and systems bus

Pentium 4 Design Reasoning

- Decodes instructions into RISC like micro-ops before L1 cache
- Micro-ops fixed length
 - Superscalar pipelining and scheduling
- Pentium instructions long & complex
- Performance improved by separating decoding from scheduling & pipelining
 - (More later – ch14)
- Data cache is write back
 - Can be configured to write through
- L1 cache controlled by 2 bits in register
 - CD = cache disable
 - NW = not write through
 - 2 instructions to invalidate (flush) cache and write back then invalidate
- L2 and L3 8-way set-associative
 - Line size 128 bytes

PowerPC Cache Organization

- 601 – 1 32kb, 8 way set associative
- 603 – 2 8kb, two way set associative
- 604 – 2 16kb, 4 way set associative
- 620 – 2 32kb, 8 way set associative
- G4
 - L1 2 32kb, 8 way set associative
 - L2 256K, 512K, 1MB

Homework

- 1. Reading book
- 2. Key Terms
- 3. Problems: 5,6,7,8,11,19